

Emergent Reasoning in Large Language Models: Mechanisms, Benchmarks, and Limitations

A Systematic Review of Chain-of-Thought and Self-Consistency Approaches

James R. Holloway

Mei-Ling Chen

Arjun Patel

Department of Computer Science, Stanford University

Allen Institute for AI, Seattle, WA

Correspondence: j.holloway@stanford.edu

Abstract

Large language models (LLMs) have demonstrated surprising emergent capabilities in multi-step reasoning tasks, often without explicit supervision on intermediate steps. This paper systematically reviews two dominant paradigms — chain-of-thought (CoT) prompting and self-consistency decoding — examining their theoretical underpinnings, empirical performance on standardized benchmarks (GSM8K, MATH, BIG-Bench Hard), and failure modes. We identify three critical limitations: sensitivity to prompt phrasing, poor calibration under distribution shift, and susceptibility to plausible-but-incorrect reasoning chains. We further propose a taxonomy of reasoning error types observed across 12 open-source and proprietary models. Our analysis suggests that current CoT approaches approximate, rather than implement, genuine symbolic reasoning, with performance gains attributable primarily to surface-level pattern matching. We conclude with directions for hybrid neuro-symbolic architectures that may address these shortcomings.

Keywords: large language models, chain-of-thought, reasoning, emergent capabilities, self-consistency, benchmarks

1. Introduction

1.1 Motivation

The rapid scaling of transformer-based language models has yielded a class of systems capable of performing tasks that were, until recently, considered exclusive to structured symbolic reasoners. Arithmetic word problems, logical deduction puzzles, and multi-hop question answering can now be addressed by models with billions of parameters using nothing more than natural language prompting. Yet the mechanisms underlying these capabilities remain poorly understood, raising fundamental questions about whether observed behaviors reflect genuine reasoning or sophisticated pattern interpolation.

1.2 Scope of This Review

We focus on published work between 2020 and 2024, covering 87 papers, 14 benchmark datasets, and empirical evaluations across 12 model families. We restrict our analysis to autoregressive transformer models and exclude retrieval-augmented systems, which introduce additional variables that complicate attribution.

2. Background and Related Work

2.1 Chain-of-Thought Prompting

Wei et al. (2022) introduced chain-of-thought prompting as a technique wherein few-shot exemplars include explicit intermediate reasoning steps, encouraging models to generate analogous reasoning chains at inference time. Subsequent work demonstrated that zero-shot CoT, achieved simply by appending 'Let's think step by step' to the prompt, yields comparable gains on many arithmetic benchmarks, suggesting that models internalize step-by-step processing schemata during pretraining. Kojima et al. (2022) provide extensive empirical support for this hypothesis across 12 diverse task types.

2.2 Self-Consistency Decoding

Wang et al. (2023) proposed self-consistency as a complementary technique: rather than relying on greedy decoding, multiple reasoning paths are sampled and the most frequently reached conclusion is selected as the final answer. This ensemble approach reduces variance attributable to single-path instability and has yielded significant gains on GSM8K (from 56.5% to 74.4% for PaLM-540B). However, the computational cost scales linearly with the number of sampled paths, limiting practical deployment.

3. Methodology

3.1 Model Selection

We evaluated GPT-4, Claude 2, LLaMA-2-70B, Mistral-7B, Falcon-40B, and seven additional models spanning parameter counts from 7B to 540B. All models were evaluated using identical prompting templates to control for prompt sensitivity. Temperature was fixed at 0.7 for sampling-based methods and 0.0 for greedy decoding baselines.

3.2 Benchmark Suite

Primary benchmarks include GSM8K (grade-school math), MATH (competition mathematics), BIG-Bench Hard (a subset of BIG-Bench tasks identified as particularly challenging for language models), and our newly introduced LogicTrace-500, a dataset of 500 first-order logic deduction problems with annotated intermediate proof steps.

4. Results and Discussion

4.1 Performance Comparison

CoT prompting consistently improved performance over standard prompting across all benchmarks, with average gains of 18.3 percentage points on GSM8K and 11.7 points on MATH.

Self-consistency provided additional improvements of 6–9 points when 40 samples were used. However, gains were substantially smaller (averaging 4.2 points) on LogicTrace-500, suggesting that formal logical deduction remains a distinct challenge from arithmetic reasoning.

4.2 Error Taxonomy

Manual analysis of 1,200 incorrect responses identified five primary error types: (i) spurious intermediate steps that do not affect the final answer but are factually incorrect; (ii) correct reasoning chains that lead to arithmetic errors in the final computation step; (iii) premise hallucination, where the model introduces unstated assumptions; (iv) logical reversal errors; and (v) context truncation in long-horizon problems. Type (iii) accounted for 31% of errors and shows no clear correlation with model scale.

4.3 Limitations

Several important limitations constrain the conclusions of this review. First, our evaluation is restricted to English-language prompts; cross-lingual generalization of CoT has been studied only superficially. Second, benchmark saturation is a growing concern: GSM8K, in particular, may have been partially memorized by models trained on data scraped post-2022. Third, human evaluation of reasoning quality is expensive and not yet standardized across the community.

5. Conclusion

This systematic review documents the impressive but bounded capabilities of chain-of-thought and self-consistency approaches in eliciting reasoning from large language models. While these techniques yield measurable improvements on standard benchmarks, the error patterns we identify suggest that current models lack robust, generalizable reasoning faculties. Future work should prioritize interpretability tools that can distinguish genuine symbolic processing from surface-level pattern completion, and the field would benefit from benchmark designs that are more resistant to data contamination. Hybrid neuro-symbolic systems that couple LLM fluency with formal inference engines represent a promising but underexplored direction.

References

- [1] Wei, J., Wang, X., Schuurmans, D., Bosma, M., Ichter, B., Xia, F., ... & Zhou, D. (2022). Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35, 24824–24837.
- [2] Kojima, T., Gu, S. S., Reid, M., Matsuo, Y., & Iwasawa, Y. (2022). Large language models are zero-shot reasoners. *Advances in Neural Information Processing Systems*, 35, 22199–22213.
- [3] Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E., Narang, S., ... & Zhou, D. (2023). Self-consistency improves chain of thought reasoning in language models. *International Conference on Learning Representations*.
- [4] Srivastava, A., et al. (2022). Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. *Transactions on Machine Learning Research*.
- [5] Hendrycks, D., et al. (2021). Measuring mathematical problem solving with the MATH dataset. *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*.

- [6] Cobbe, K., et al. (2021). Training verifiers to solve math word problems. arXiv preprint arXiv:2110.14168.
- [7] Lightman, H., et al. (2023). Let's verify step by step. arXiv preprint arXiv:2305.20050.